

CS475 Spring 2021 Homework 3

Team 0: Dongkwan Kim, Jiseon Kim, and Alice Oh

April 15, 2021

BERT [1] has become a standard architecture for NLP research ever since it was published. BERT computes the representation for every token, but for sentence-level tasks (e.g., sentiment analysis), we usually use the output representation of the special token [CLS]. Is this the best way to pool token-level representations to create the sentence-level representation?

In this homework, your team will implement and compare different techniques for pooling token-level representation in BERT, namely **BERTPooler** in **huggingface** library (<https://huggingface.co/>).

1 Tasks

One simple pooling method is to apply permutation-invariant operations (e.g., mean, max, sum) across all token representations. We define the output representation of **MeanMaxTokens** (shortly MMT) as:

$$C_{\text{MMT}} = \left[\sum_{i=0}^{L-1} T_i \parallel \max_i T_i \right], \quad C_{\text{MMT}} \in \mathbb{R}^{2H}, \quad T_i \in \mathbb{R}^H, \quad (1)$$

where L is the maximum length of sequence and \parallel is the vector concatenation operation. Note that we borrow notations from the original BERT paper [1].

The **BERTPooler** in **huggingface** implementation applies linear transformation with $W \in \mathbb{R}^{H \times H}$ and tanh activation (See **BertPooler** class). Similarly, we apply linear transformation with $W_{\text{MMT}} \in \mathbb{R}^{H \times 2H}$ and tanh activation.

$$C = \tanh(C_{\text{MMT}} W_{\text{MMT}}^T), \quad C \in \mathbb{R}^H \quad (2)$$

The tasks in this homework are as follows:

1. Implement the **MeanMaxTokens** pooler (See **MeanMaxTokensBertPooler** class in **bert_poolers.py**).
2. Implement your own BERT pooler (See **MyBertPooler** class) and describe its architecture and rationale in your report. It does not have to be completely novel.
3. Choose one dataset in GLUE [2], and compare the test performance of three poolers (See **run_glue.py**).
4. Discuss the result. Negative results are fine, the point is how you interpret and explain it.

2 Submission

The files you should submit are

1. Your team's **bert_poolers_{team_no}.py** (e.g., **bert_poolers_0.py**).
2. Your team's two-page **report_{team_no}.pdf** (e.g., **report_0.pdf**). Use this L^AT_EX file as a template, and do not change style attributes in this file. References are not included in the page-limit.

3 Grading

Comprehensive evaluation based on clarity, validity, and interestingness. You will get zero points if you violate academic integrity (e.g., plagiarism and data manipulation).

References

- [1] Jacob Devlin, Ming-Wei Chang, Kenton Lee, and Kristina Toutanova. BERT: Pre-training of deep bidirectional transformers for language understanding. In *Proceedings of the 2019 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, Volume 1 (Long and Short Papers)*, pages 4171–4186, Minneapolis, Minnesota, June 2019. Association for Computational Linguistics.
- [2] Alex Wang, Amanpreet Singh, Julian Michael, Felix Hill, Omer Levy, and Samuel R. Bowman. GLUE: A multi-task benchmark and analysis platform for natural language understanding. In *International Conference on Learning Representations*, 2019.